


PRODUCT / GP SCORING AND BENCHMARKING

Every fund placed against peers and public markets.

Alts Cortex is an AI research platform for private markets. One selected manager (GP) and fund stay in view across returns, cash flows, holdings, benchmarks and reports. Python computes each figure; AI writes the explanation.

PEER DISCIPLINE



Vintage stays fixed.

A peer cell holds funds of one vintage year, strategy and region. A small cell widens geography, then strategy, then asset class.

ANNUALISED RETURN, 13 REAL FUNDS



FIVE MEASURES OF ONE FUND

Peer quartiles

Net IRR, total value, paid-out and held-value multiples.

Public-market equivalent

The fund against an index bought on the same dates.

Public comparables

Sector multiples for 2026 from 3,208 listed companies.

Manager score

A score from 0 to 100 on total value and cash returned.

League table

Managers ranked by the average quartile of their funds.

BENCHMARK SET

5,076

Peer cells simulated universe

22,383

Managers ranked simulated, by quartile

58

Sector indices daily, some from 1980

14


Real funds from public filings

Source: System Overview / Managers and funds, 7 Oct 2026. Peer cells and the manager league are a simulated vendor universe, labelled on every screen; 14 funds come from public filings.

PRODUCT / SECONDARY PRICING

From tender offer to priced verdict.

An annual cash-flow model pays for the stake, funds the unfunded commitment, collects distributions and values the exit. The pricing model sets the verdict; the desk decides the bid.



FOUR PRICES, FOUR PURPOSES

01	Market indication Median of comparable successful bids, or a range under five comparables.
02	Underwritten value Reference NAV moved by sector exposure since its date.
03	Walk-away price The highest price that still earns an assumed 15 percent a year.
04	Proposed bid The lower of the walk-away price and the comparable median; the desk decides.

STATED STRESS CASES

15% Base hurdle assumed annual return	18% Rates up NAV 10% lower	+12 mo Exits slow exit one year later
--	---	--

PORTFOLIO FIT

16 limits, 6 filed and 10 assumed, cover strategy, region, deal type, vintage, manager, borrowing and liquidity. Each limit is rechecked before and after the deal.


41 deal reports	27 bid with conditions and 14 pass, written from 1,815 tagged facts.
------------------------	--

Source: System Overview / Secondaries, 7 Oct 2026. Hurdles and assumed limits are stated model settings, not universal buyer rules. 11 transaction frameworks run from LP stakes to structured deals.

SYSTEM / EXTRACTION AND QUALITY CONTROL

Two AI readers. Python checks during reading.

Each PDF page reaches two independent AI readers as a picture, page-marked text and a number grid. Fixed Python checks hold each reader to the printed page, and a third reader rules every difference.



93.3%

Two-reader agreement

Both readers typed the same value for 7,477 of 8,016 shared fields in 53 PDF reports. The third reader ruled the other 539.

930 documents read: 53 PDF reports and 877 SEC filings.

PYTHON CHECKS DURING READING

Table check	A table page with zero rows needs a reason.
Keyword check	An amount with zero rows needs a reason.
Excuses	Reading difficulty is refused as a reason.
Quote check	Each value keeps a quote from its page.
Early check	Each reader validates its first three pages.

CHECKS AFTER READING

Match	Both readers' rows are paired by cited line.
Compare	Equal, value conflict, label conflict or one-sided.
Split	Equal values stay; differences go to the third reader.
Rule	The third reader accepts A or B, merges, adds or rejects.
Final check	The finished file is validated again.

Source: System Overview / Data, 7 Oct 2026. Agreement covers shared fields in the measured sample and is distinct from an end-to-end error rate. One reader cost about \$0.034 per page on 3 measured reports.

SYSTEM / CORTEX AND TEMPORAL

Cortex proves the links. Temporal repairs the gaps.

Cortex is a proof graph: each line between managers, funds, holdings and filings has a receipt naming its source rows and method. Temporal runs the recorded investigations that rule on gaps and conflicts in the data.

CORTEX / PROOF GRAPH

11.9 million

Certified lines, 0 blocked

403,916 facts: 58 from real filings (33 funds, 25 companies) and 403,858 simulated vendor records, labelled as simulated.

	Search and RAG	Cortex
Finds	text that sounds like the question	connections the records prove
Evidence	a closeness score only	each line graded by its evidence
Answer check	left to the reader	each figure matched to its record

TEMPORAL / RECORDED INVESTIGATIONS

- **Detect**
A scan after each release lists blank values, unreconciled totals and filings that print different values.
- **Investigate**
Cortex suggests a likely cause with its evidence. Two AI agents re-read the source documents.
- **Rule**
A third agent rules. Each ruling is logged with its evidence: 62 so far. Temporal keeps the run history.
- **Release**
A repair reaches the database through an approved release. Each job resumes after a crash at the step it stopped on.

Cortex and Temporal together make the record self-correcting: bad data is found, investigated, ruled on and repaired.


Source: System Overview / Cortex and Data, 7 Oct 2026. A receipt supports its stated relationship only. Repairs reach the database through an approved release.

AI / REPORTS, CHECKS AND TRACEABILITY

Code computes. AI explains.
Every number is checked.

DUE DILIGENCE ON ONE FUND

AMG Pantheon Master Fund



6 Answered 13 Answered in part 3 Open

22 default questions for this fund: evidence coverage, not a score of manager quality.

- Answered Every required fact and document is on file.
- In part Some facts are on file; one item stays open.
- Open All needed facts remain to be found.

Each open item becomes a request naming the document and its holder.

THE NUMBER CHECK

A sentence with an unsupported
number is rejected.

Each number in a draft is matched to the record. A failing sentence is removed and listed with its reason. Interpretation stays marked as unverified.

TAGS IN EVERY REPORT

- F01 A fact from a filing. One click opens the filing at the cited place.
- Q1 A fact from outside research. One click opens that source.
- ASSUMPTION A buyer input or model setting, labelled beside the figure.

Origins: Reported • Derived • Estimated • Simulated

Five report types share one method: risk profile, new analysis, due diligence, deal and industry reports. Python computes the figures and the recommendation, the model writes the words, and a fixed check keeps only sentences with numbers found in the record.

SKILLS AND SCHEMAS THE AGENTS LOAD

- 8 Written specialist skills loaded by each fund risk report
- 258 Diligence checks behind the due diligence skill
- 1,300 Research passages 109 files; 179 library reports
- 7 Schema sections the fixed outline of a deal report
- 33 Schema questions due diligence, in 11 families

Every data point verifiable in one click.

Every AI number checked by code.

Ready for the investment committee.

Source: System Overview / Due diligence reports and Managers and funds, 7 Oct 2026. 9,092 of 10,652 checked references open at the cited place; the rest open the filing.